

FULL SEO AUDIT REPORT

betteragent.dev

Prepared by Claude SEO

38/100

SEO Health Score

June 18, 2026 Full Audit

Table of Contents

1. Executive Summary

Key Metrics, Critical Issues & Quick Wins

2. Audit Categories

What Works

Findings by Severity

3. Action Plan

Phased Roadmap

4. Data Sources & Methodology

1. Executive Summary

This report presents a comprehensive SEO audit of **betteragent.dev** for a **B2B SaaS / developer tool (AI agent infrastructure for SaaS products)** site, generated on June 18, 2026. Findings combine technical, content, schema, performance, visual, and search-readiness evidence as available.

38/100

SEO HEALTH SCORE

Critical Issues Found:

1. **Info:** No structured data (JSON-LD) anywhere on either domain
2. **Info:** Six+ distinct pages (/ , /pricing, /demo, /contact, /auth/sign-in, /auth/sign-up) share the exact same <title> and meta description
3. **Info:** No sitemap.xml on betteragent.dev or docs.betteragent.dev (both 404)
4. **Info:** robots.txt blocks GPTBot, ClaudeBot, Google-Extended, Applebot-Extended, CCBot, Bytespider, meta-externalagent
5. **Info:** /demo is fully client-rendered with no SSR fallback -- 9 words of server text, no <h1>, duplicate meta

Quick Wins:

1. Write unique titles/descriptions per page (~10 min/page, highest leverage available)
2. Add app/sitemap.ts and reference it in robots.txt
3. Add Open Graph + Twitter Card tags with one branded image
4. Add canonical tags via Next.js metadata API (alternates.canonical)
5. Add Organization + SoftwareApplication JSON-LD to the homepage using existing pricing data

2. Audit Categories

2.1 Technical SEO

Score: **55/100**

What Works

- HTTPS enforced with HSTS (max-age=63072000)
- Clean, semantic URLs (/pricing, /contact, /demo, /privacy, /terms)
- robots.txt explicitly allows general search crawling (Content-Signal: search=yes)
- Mobile viewport meta present on every page
- Pages are Next.js-prerendered and edge-cached (x-vercel-cache: HIT, cf-cache-status)
- Custom 404 page returns proper 404 status code with helpful copy and links

Findings

No sitemap.xml on either domain **Critical**

<https://betteragent.dev/sitemap.xml> and <https://docs.betteragent.dev/sitemap.xml> both resolve to the site's not-found page (soft 404, HTTP 404 confirmed). No app/sitemap.ts exists in the codebase. No Sitemap: directive in either robots.txt.

Recommendation: Add app/sitemap.ts exporting a MetadataRoute.Sitemap with all public routes (/ , /pricing, /demo, /contact, /privacy, /terms); add an equivalent for the Nextra docs app; reference both via Sitemap: lines in robots.txt.

No canonical tags on any page **High**

None of the 8 crawled pages (/ , /pricing, /demo, /contact, /privacy, /terms, /auth/sign-in, /auth/sign-up) emit a <link rel="canonical">. This matters more here because the apex domain redirects to www.betteragent.dev, and metadataBase in app/layout.tsx is set to the apex (<https://betteragent.dev>) -- a mismatch with the domain actually serving content.

Recommendation: Set alternates.canonical per page via the Next.js Metadata API, and align metadataBase with whichever host (apex or www) is the canonical one.

robots.txt blocks every major AI crawler **High**

Cloudflare's managed robots.txt content disallows GPTBot, ClaudeBot, Google-Extended, Applebot-Extended, CCBot, Bytespider, meta-externalagent, Amazonbot, and CloudflareBrowserRenderingCrawler, with Content-

Signal: ai-train=no set globally. This looks like Cloudflare's default AI-bot ruleset rather than a deliberate product decision.

Recommendation: Confirm with the team whether this is intentional. If the goal is just to prevent training on the content, ai-train=no can stay, but the explicit Disallow: / rules for ClaudeBot/GPTBot/Google-Extended also block ai-input (RAG/grounding) -- meaning the product can never surface when someone asks an AI assistant 'what should I use to add an agent to my SaaS.' Loosen to allow read/input access while keeping training disabled.

No baseline security headers **Medium**

The homepage response has no Content-Security-Policy, X-Content-Type-Options, X-Frame-Options, Referrer-Policy, or Permissions-Policy headers.

Recommendation: Add these via next.config headers() or a middleware -- low effort, and increasingly factored into some search trust/quality signals.

Apex-to-www redirect uses a temporary (307) status **Low**

https://betteragent.dev/ responds with HTTP 307 to https://www.betteragent.dev/ instead of a permanent redirect.

Recommendation: Switch to a 301/308 permanent redirect so link equity consolidates onto the canonical host and crawlers don't keep re-checking the apex.

2.2 Content Quality

Score: **40/100**

What Works

- Homepage copy is specific and avoids generic SaaS filler ('it reads your routes and server actions, generates the schemas...')
- Pricing page clearly explains the credit-based billing model with a full plan comparison
- Docs subdomain (Nextra) has clean, purpose-built MDX content: quickstart, AI setup, CLI reference, tool files, components

Findings

/demo renders almost no content server-side **Critical**

/demo's server-rendered HTML contains only ~9 words of text (the page title, duplicated from the homepage) versus ~1,000 words on the homepage. The page is a bare `<Suspense><VariantSwitcher /></Suspense>` client component with no metadata export, no `<h1>`, and no static fallback content.

Recommendation: Add a server-rendered heading, short description, and static preview/screenshot as a Suspense fallback so the page has indexable content even before the client-side variant switcher hydrates.

No authority/E-E-A-T signal anywhere in the site **High**

No About, team, founder, case-study, or testimonial page exists in the link graph. There's no human-authorship or credibility signal beyond the product copy itself.

Recommendation: Add a short About/team page, and surface at least one customer quote or usage stat once available.

No blog or long-tail educational content **Medium**

/blog 404s; there is no content marketing surface beyond the docs.

Recommendation: Not urgent pre-PMF, but flag as a growth lever -- the docs subdomain is already structured well enough to extend with guides.

2.3 On-Page SEO

Score: **28/100**

What Works

- Single `<h1>` present on most pages (`/`, `/pricing`, `/contact`)
- `lang="en"` correctly set on `<html>`
- Heading hierarchy is mostly clean on the homepage (one H1, H2 sections, H3 subsections)
- Internal navigation links to all key public routes (pricing, demo, contact, privacy, terms, sign-in, sign-up) plus the docs subdomain

Findings

Six pages share an identical title and meta description **Critical**

`/`, `/pricing`, `/demo`, `/contact`, `/auth/sign-in`, and `/auth/sign-up` all serve the exact same `<title>BetterAgent` -- the agent layer your SaaS is missing `</title>` and identical meta description, because only `app/layout.tsx` defines metadata and none of these `page.tsx` files override it.

Recommendation: Add a page-level ``export const metadata`` (or `generateMetadata`) to each of `pricing/page.tsx`, `demo/page.tsx`, `contact/page.tsx`, and the two auth pages with a unique, page-specific title and description.

Doubled brand suffix on /privacy and /terms titles **Medium**

Both pages set their own title to '...-- BetterAgent' while the root layout's title.template ('%s -- BetterAgent') also appends the suffix, producing 'Privacy Policy -- BetterAgent -- BetterAgent' and 'Terms of Service -- BetterAgent -- BetterAgent' in the rendered <title>.

Recommendation: Drop the manual '-- BetterAgent' suffix from these two pages' title strings and let the layout template add it once.

Zero Open Graph or Twitter Card tags **High**

No og:title, og:description, og:image, og:url, twitter:card, or related tags exist on any crawled page. Links shared in Slack, X, or LinkedIn render with no title, description, or image.

Recommendation: Add openGraph and twitter fields to the Metadata API (can be set once at the root and overridden per page) plus a branded 1200x630 opengraph-image asset.

/demo has no <h1> **Medium**

Same root cause as the content-quality finding -- the page has no server-rendered heading at all.

Recommendation: Covered by the /demo SSR fix above.

2.4 Schema / Structured Data

Score: **0/100**

Findings

No JSON-LD anywhere on either domain **High**

No application/ld+json block was found on any of the 8 crawled betteragent.dev pages or on docs.betteragent.dev. No Organization, SoftwareApplication, Product/Offer, or BreadcrumbList schema exists.

Recommendation: Add Organization schema sitewide (name, logo, url, sameAs) and SoftwareApplication or Product+Offer schema on /pricing, mapped directly from the existing plan comparison table -- this is a near-zero-content-change, high-leverage fix.

2.5 Performance (CWV)

Score: **65/100**

What Works

- Prerendered Next.js output served via Vercel with edge caching (x-vercel-cache: HIT, x-nextjs-prerender: 1)
- Cloudflare CDN in front with HTTP/3 (alt-svc h3) and gzip compression
- Homepage TTFB ~260ms and full response in ~270ms from a cold curl, consistent with edge-cached delivery
- Reasonable HTML payload (~108KB uncompressed) for a content-rich landing page

Findings

No lab-measured Core Web Vitals available **Info**

PageSpeed Insights/CrUX API access is not configured in this environment, so LCP/INP/CLS are not directly measured here. The score above is a heuristic based on response headers, caching behavior, and architecture, not a Lighthouse run.

Recommendation: Run PageSpeed Insights manually (or configure a `GOOGLE_API_KEY` for python scripts/`google_report.py`) to get real LCP/INP/CLS numbers, especially given the homepage loads 16+ separate JS chunks.

2.6 AI Search Readiness

Score: **20/100**

What Works

- Docs subdomain content (quickstart, AI setup, CLI reference, tool files, components) is clean, well-segmented MDX -- good raw material for citation once accessible
- Page copy itself is concrete and specific rather than vague marketing language, which tends to extract well into AI answers when crawlable

Findings

Every major AI assistant/answer-engine crawler is blocked **Critical**

GPTBot, ClaudeBot, Google-Extended, Applebot-Extended, CCBot, Bytespider, and meta-externalagent are all disallowed in robots.txt. For a product whose entire pitch is 'the agent layer your SaaS is missing,' this means it cannot surface when a developer asks ChatGPT, Claude, or Gemini what to use to add an agent to their product.

Recommendation: Allow ai-input access for these crawlers (RAG/grounding at query time) even if ai-train=no is kept for IP reasons -- the two are independently controllable via Content-Signal.

No llms.txt on either domain **Medium**

Neither `betteragent.dev` nor `docs.betteragent.dev` serves an `llms.txt` file (both 404).

Recommendation: Add an llms.txt at the docs subdomain root pointing at the existing quickstart/ai-setup/cli/tools/components pages -- low effort given the content already exists.

2.7 Images

Score: **50/100**

What Works

- No broken or oversized raster images -- the marketing pages are effectively zero-- A 512x512 icon.png is correctly referenced via <link rel="icon">

Findings

No og:image / social preview image **High**

No opengraph-image route or static asset exists -- confirmed via /opengraph-image returning 404. Shared links have no visual preview.

Recommendation: Add app/opengraph-image.tsx (or a static 1200x630 PNG) using Next.js's built-in OG image generation.

No favicon.ico, apple-touch-icon, or manifest.json **Low**

/favicon.ico, /apple-touch-icon.png, and /manifest.json all 404 (after following the apex->www redirect). Only the 512x512 icon.png is served.

Recommendation: Add a favicon.ico fallback for older clients, an apple-touch-icon for iOS home-screen bookmarking, and a basic manifest.json if any PWA-like behavior is desired.

3. Action Plan

Phase 1: Critical Fixes Week 1

- Write unique <title> + meta description for /pricing, /demo, /contact, /auth/sign-in, /auth/sign-up (currently all identical to the homepage)
- Add app/sitemap.ts for betteragent.dev and an equivalent for docs.betteragent.dev; add Sitemap: lines to both robots.txt files
- Server-render real content (heading, description, static preview) on /demo -- currently ~9 words of text, fully client-only
- Decide and fix the AI-crawler block in robots.txt -- allow ai-input for ClaudeBot/GPTBot/Google-Extended at minimum, even if ai-train=no is kept
- Fix the doubled '-- BetterAgent -- BetterAgent' title suffix on /privacy and /terms

Phase 2: High-Impact Improvements Weeks 2-3

- Add canonical tags sitewide via Next.js alternates.canonical, and align metadataBase with the actual canonical host
- Add Open Graph + Twitter Card tags with one branded 1200x630 image (none exist today)
- Add Organization + SoftwareApplication/Product JSON-LD -- map directly from the existing /pricing comparison table
- Add llms.txt pointing into docs.betteragent.dev (quickstart, ai-setup, cli, tools, components)

Phase 3: Content & Authority Month 2

- Add an About/company page (founder/team info -- zero authority signal exists today)
- Add at least one customer testimonial or usage stat to the homepage
- Switch apex->www redirect from 307 to 301/308
- Add baseline security headers (CSP, X-Content-Type-Options, Referrer-Policy, Permissions-Policy)
- Add favicon.ico, apple-touch-icon, and manifest.json

Phase 4: Monitoring & Iteration Ongoing

- Run PageSpeed Insights/CrUX with a Google API key for real Core Web Vitals (current Performance score is a heuristic, not measured)
- Consider extending docs with guides/blog content for long-tail and AI-citation surface area once core priorities allow

Data Sources & Methodology

Source	Description	Update Frequency
PageSpeed Insights API	Lighthouse lab audit (mobile emulation, Moto G Power, slow 4G)	Real-time
Chrome UX Report (CrUX)	28-day rolling field data from real Chrome users	Daily ~04:00 UTC
CrUX History API	25-week p75 trend data per metric	Weekly
Google Search Console	Search Analytics (clicks, impressions, CTR, position)	2-3 day lag
URL Inspection API	Per-URL index status, coverage state, crawl info	Real-time (2,000/day)

Report generated by Claude SEO — Google SEO Intelligence Skill — June 18, 2026
Methodology based on Google Web Vitals thresholds, Search Console documentation, and Lighthouse scoring algorithms.